



# CS 4200 Final Project

Aryan Miriyala

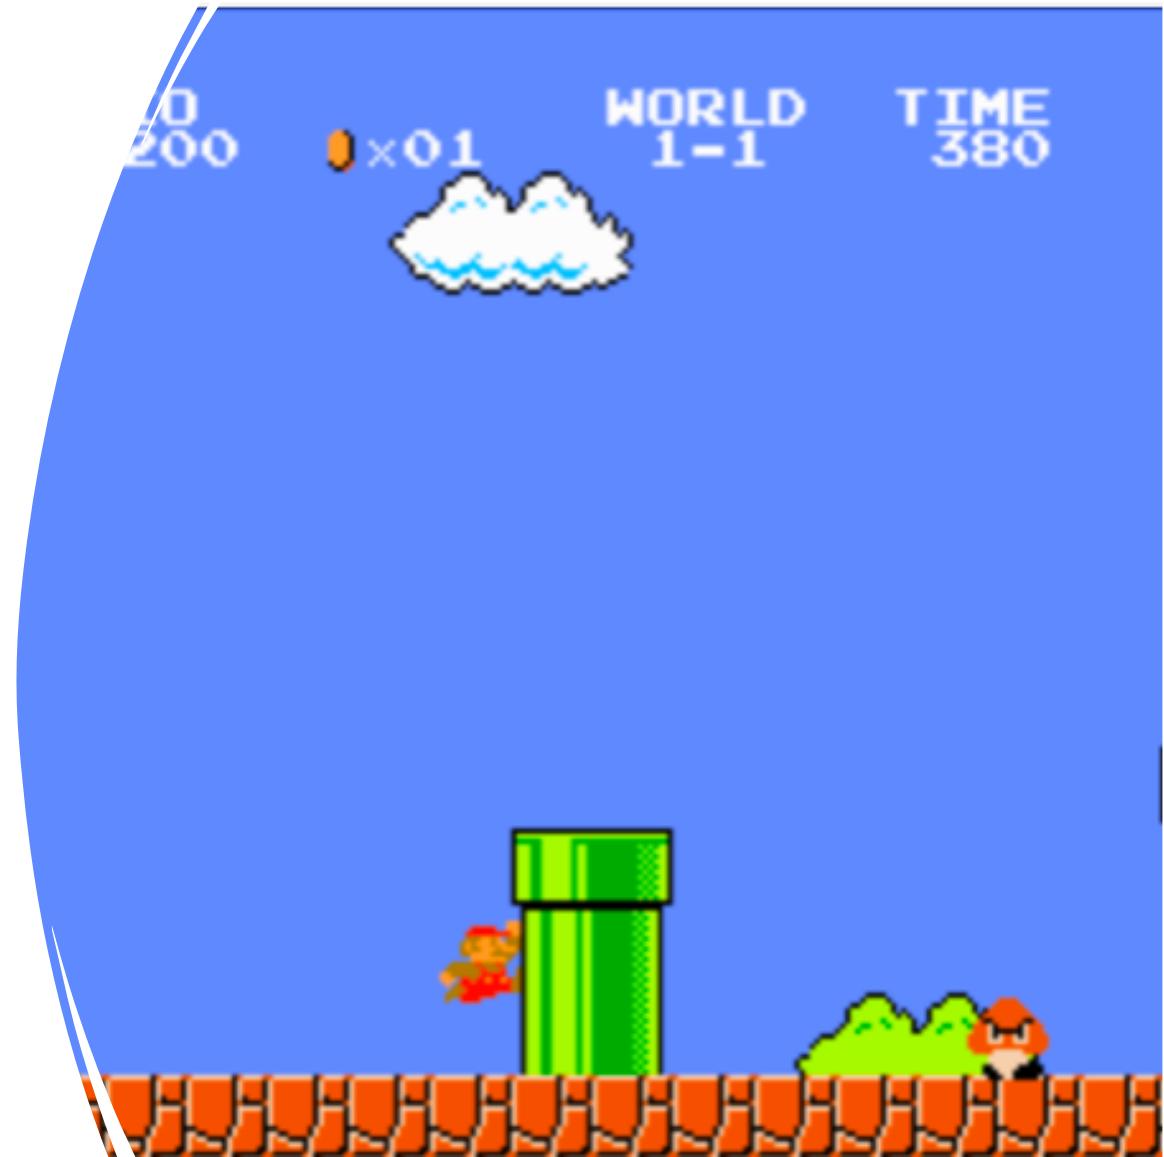
Cedric Dortch

Super Mario Reinforcement Learning Agent

Group 16

# Contents in this Presentation

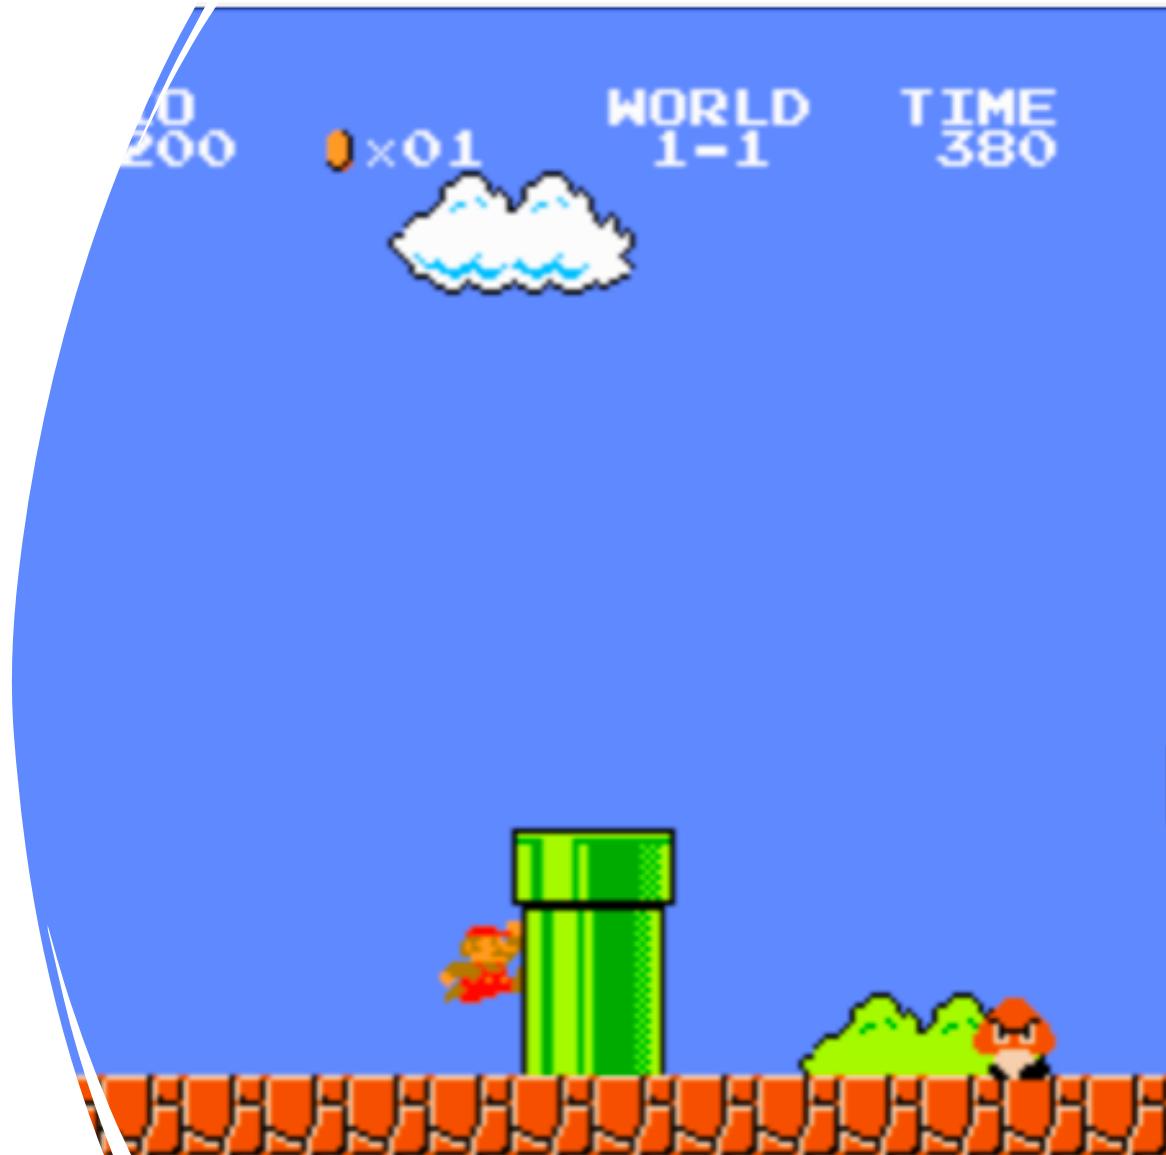
- Why Super Mario?
- Defining the Environment
- Reinforcement Learning
- Proximal Policy Optimization
- Deep Q-Network
- Results
- Challenges
- Conclusion



# Why Super Mario?

---

- Popular game and an interesting topic in the context of AI
- Dynamic environment with tons of strategies
- The game is easily accessible through OpenAI Gym
- Extensive studies have been made on Super Mario Ai



# Defining the environment

- Super Mario World PEAS representation
  - Performance: Reach the end of the level
  - Environment: The game world, the enemies, platforms, power ups
  - Actuators: The actions that are taken by Mario. We are restricted to SIMPLE\_MOVEMENTS this time.
  - Sensors: Elements that provide information about the current state, position of Mario, obstacles and enemies.
- Action Space is restricted to cut down the amount of data that needs to be consumed
- A technique called gray scaling is used to cut down the data to be processed to 1/3rd

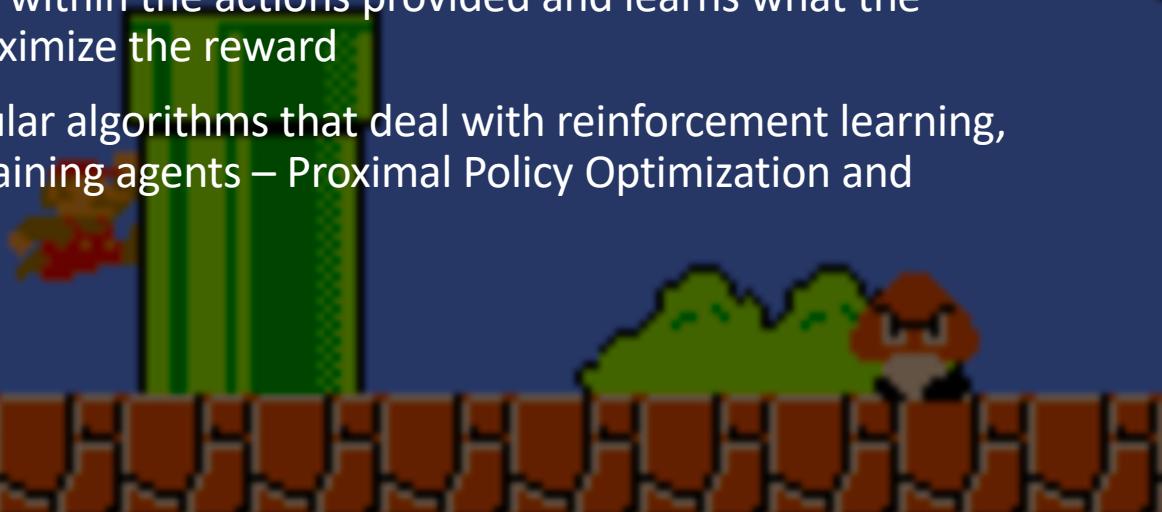
# Defining the environment

- We utilized the concept of wrapping our environment in different contexts to make use of several tricks to train our model
- Super Mario has a huge state space and action space, so brute force would take forever to converge
- Learning rate: Telling our model how important the new information learned should be. We take a really low number for a gradual convergence to the optimal value
- Gray Scaling, Frame Stacking, Resize Wrapping and Skip Framing are some concepts we are going to employ

# Reinforcement Learning – Our Approach

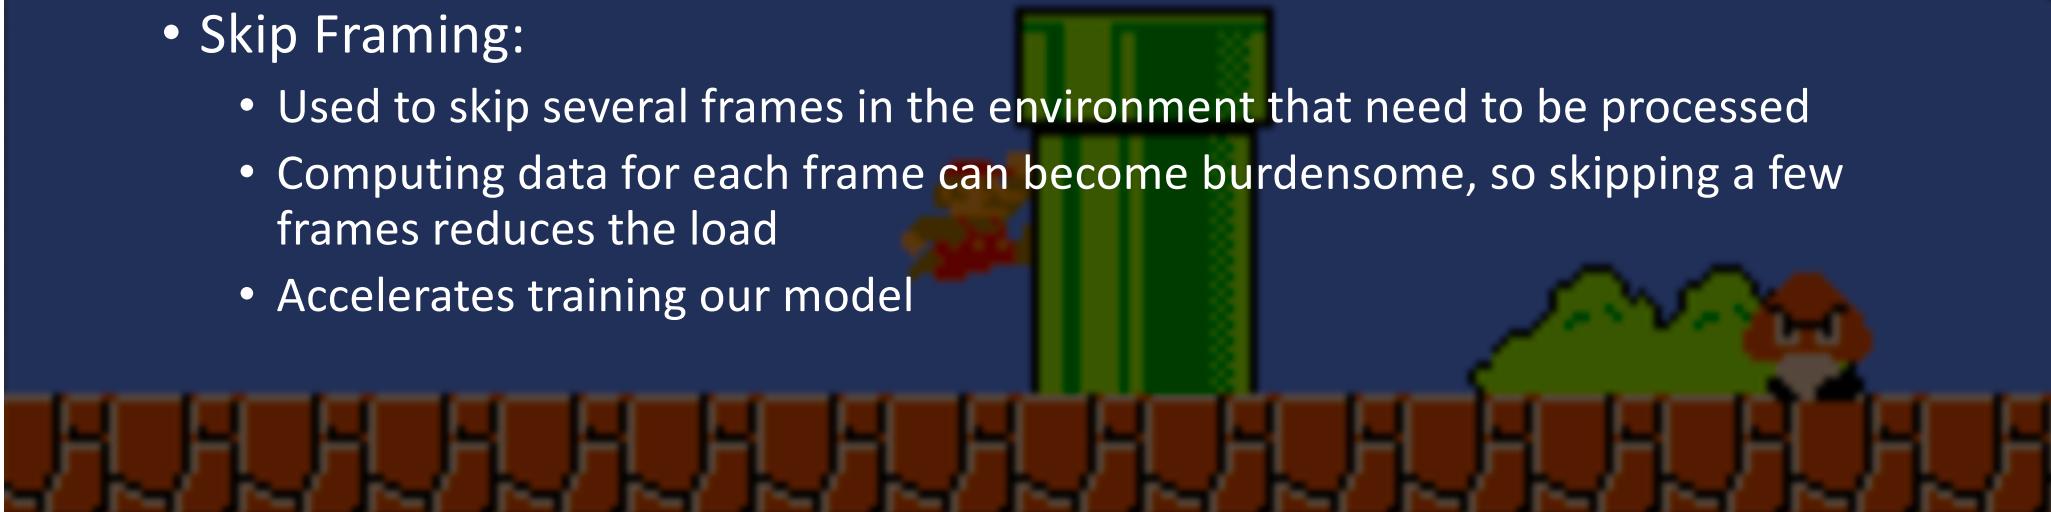
---

- Learning how to maximize expected rewards based on observed transitions
- A machine learning approach to help software make decisions
- Imagine a Markov Decision Process, but we have no idea what the transition probabilities and reward functions are
- The agent explores the environment within the actions provided and learns what the optimal actions for a state are to maximize the reward
- We are going to talk about two popular algorithms that deal with reinforcement learning, especially when it comes to game training agents – Proximal Policy Optimization and Deep Q-Network



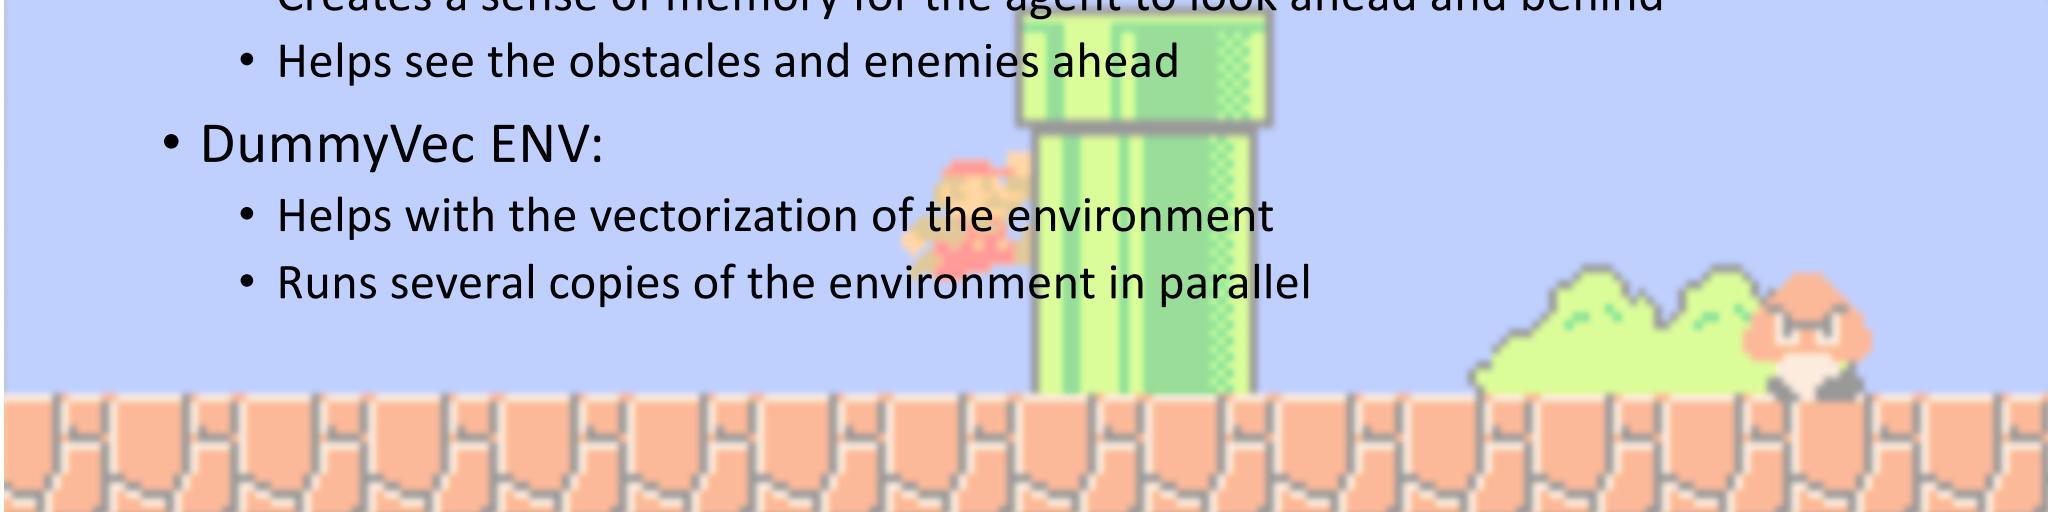
# Let's talk about a few strategies before

- Gray Scaling:
  - A strategy that converts colored images to gray scales, which in turn reduces the amount of computation we need to do and store
  - RGB colors are converted to gray scale which reduces the amount of computation to a third
- Skip Framing:
  - Used to skip several frames in the environment that need to be processed
  - Computing data for each frame can become burdensome, so skipping a few frames reduces the load
  - Accelerates training our model



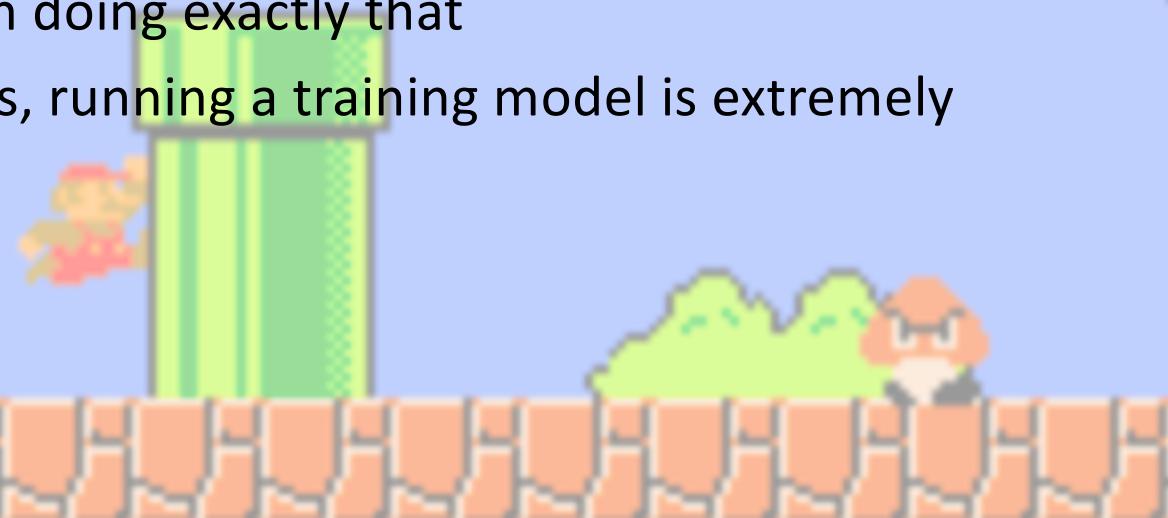
# Let's talk about a few strategies before

- Resize Operation:
  - Resize the resolution of the frames processed to reduce data size
- Frame Stacking:
  - Stacks multiple frames together for multi framed observation
  - Creates a sense of memory for the agent to look ahead and behind
  - Helps see the obstacles and enemies ahead
- DummyVec ENV:
  - Helps with the vectorization of the environment
  - Runs several copies of the environment in parallel



# Why talk about these concepts?

- The computation we will be performing to achieve a model close to optimal convergence is very high
- Only way to do it with the resources available is to utilize strategies that reduce the amount of data that we look at
- These strategies help us with doing exactly that
- Even with all these strategies, running a training model is extremely difficult
- We will talk about this later



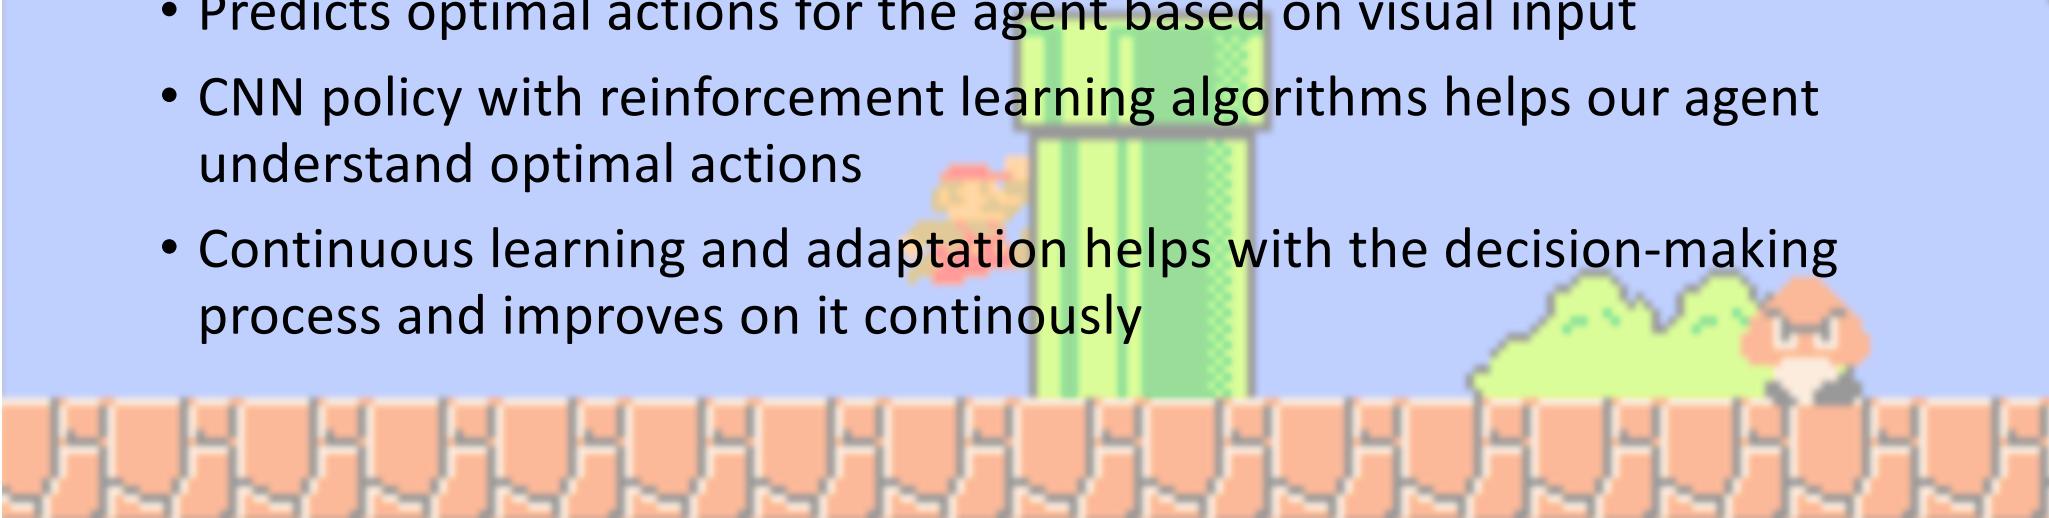
# CNN Policy (Convolved Neural Network)

- A specialized neural network architecture for analyzing visual data
- Utilized convoluted layers to detect features such as the environment that it is in, the hurdles and the enemies
- The policy automatically extracts from raw visual data
- Very helpful for interpreting image frames, such as when gaming
- Helps decision making for agents that are navigating complex environments



# CNN Policy for a Mario agent

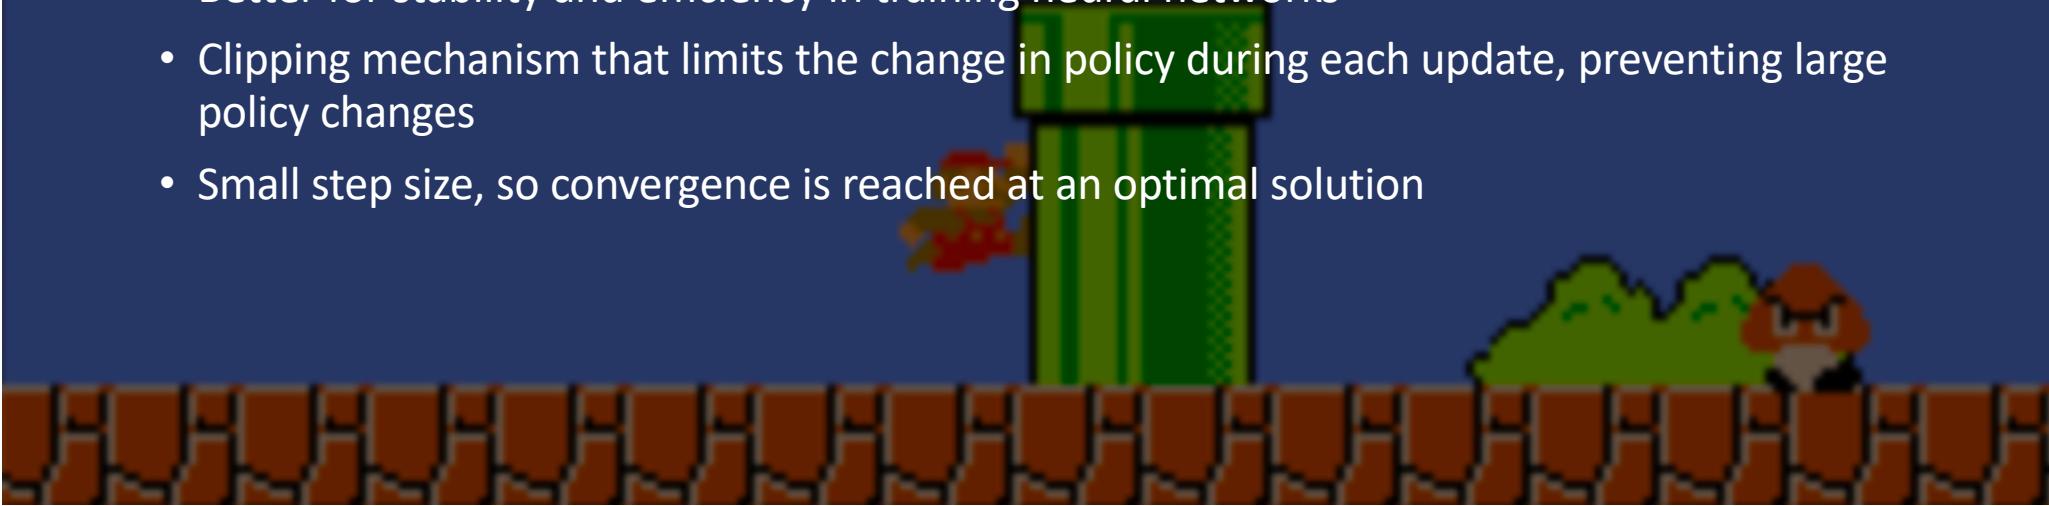
- We are going to process the environment as a stack of frames or images that contain raw data for the agent
- CNN policy analyzes game screen images to recognize important elements such as Mario, enemies, obstacles and power ups
- Predicts optimal actions for the agent based on visual input
- CNN policy with reinforcement learning algorithms helps our agent understand optimal actions
- Continuous learning and adaptation helps with the decision-making process and improves on it continuously



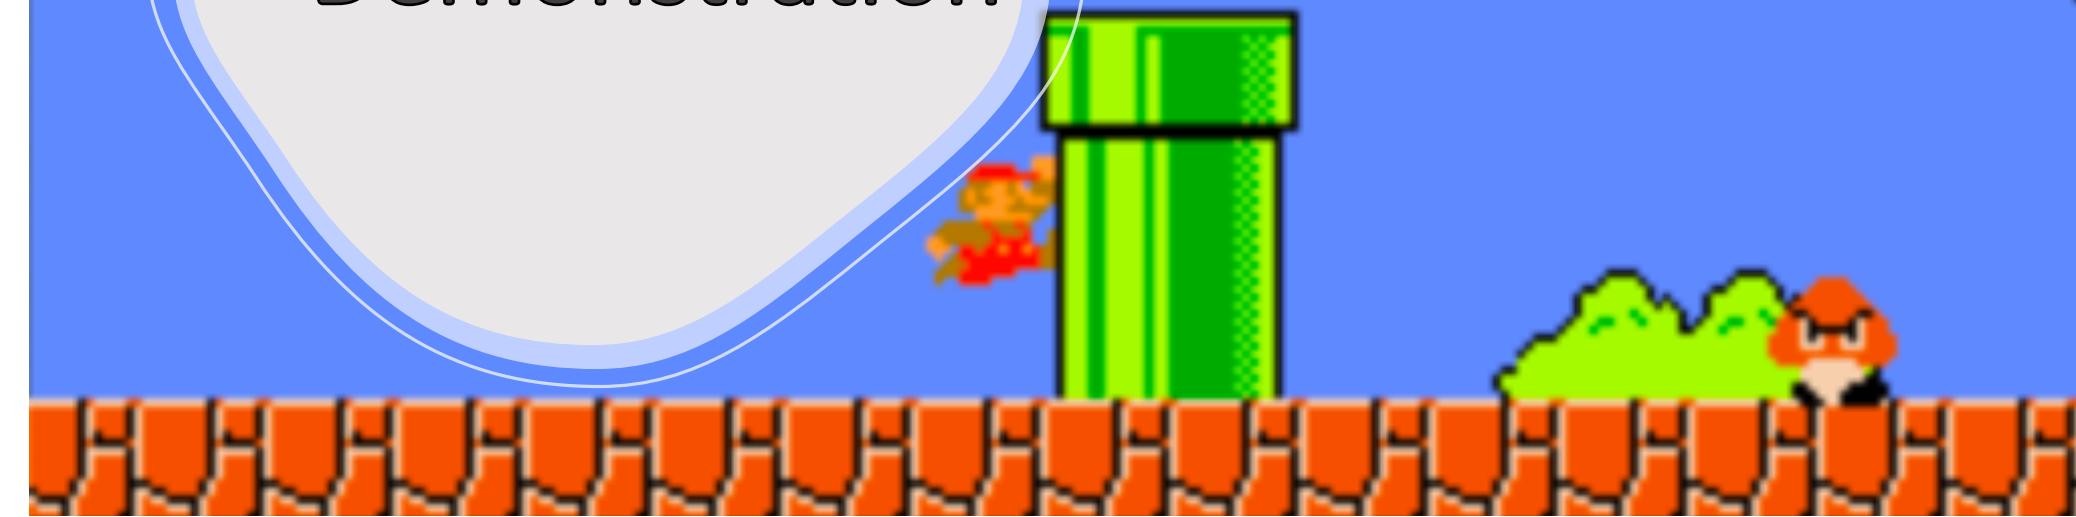
# Proximal Policy Optimization (PPO)

---

- Reinforcement Learning algorithm that trains agents to perform tasks where the actions impact future states and rewards
- Focuses on optimizing the policy function mapping observed states to relevant actions
- Iterative updates help the agent learn to make better decisions, optimizing over time
- Better for stability and efficiency in training neural networks
- Clipping mechanism that limits the change in policy during each update, preventing large policy changes
- Small step size, so convergence is reached at an optimal solution



# Code Demonstration



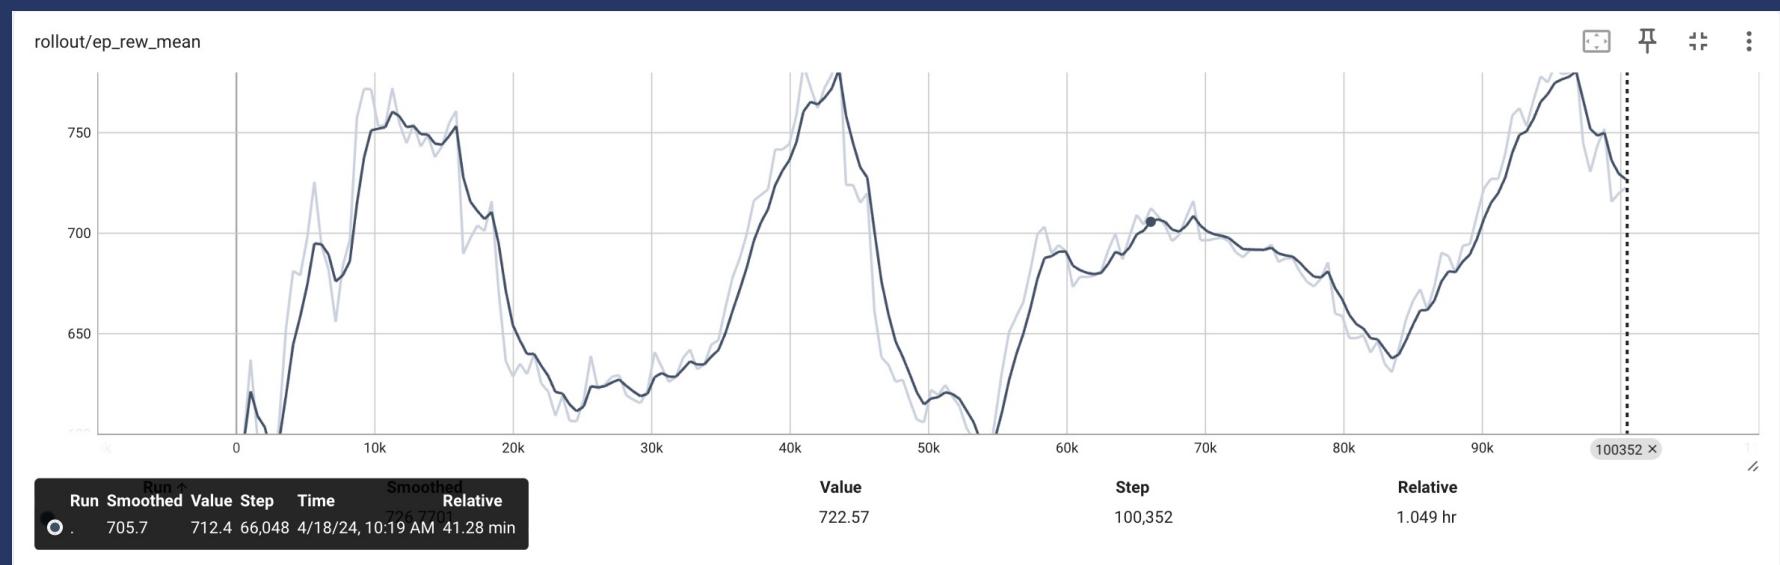
# Running the Training Model

|                      |           |
|----------------------|-----------|
| rollout/             |           |
| ep_len_mean          | 27        |
| ep_rew_mean          | 252       |
| time/                |           |
| fps                  | 24        |
| iterations           | 553       |
| time_elapsed         | 11634     |
| total_timesteps      | 283136    |
| train/               |           |
| approx_kl            | 0.0       |
| clip_fraction        | 0         |
| clip_range           | 0.2       |
| entropy_loss         | -5.05e-06 |
| explained_variance   | 1         |
| learning_rate        | 0.001     |
| loss                 | 0.0414    |
| n_updates            | 5520      |
| policy_gradient_loss | 1.03e-07  |
| value_loss           | 0.131     |

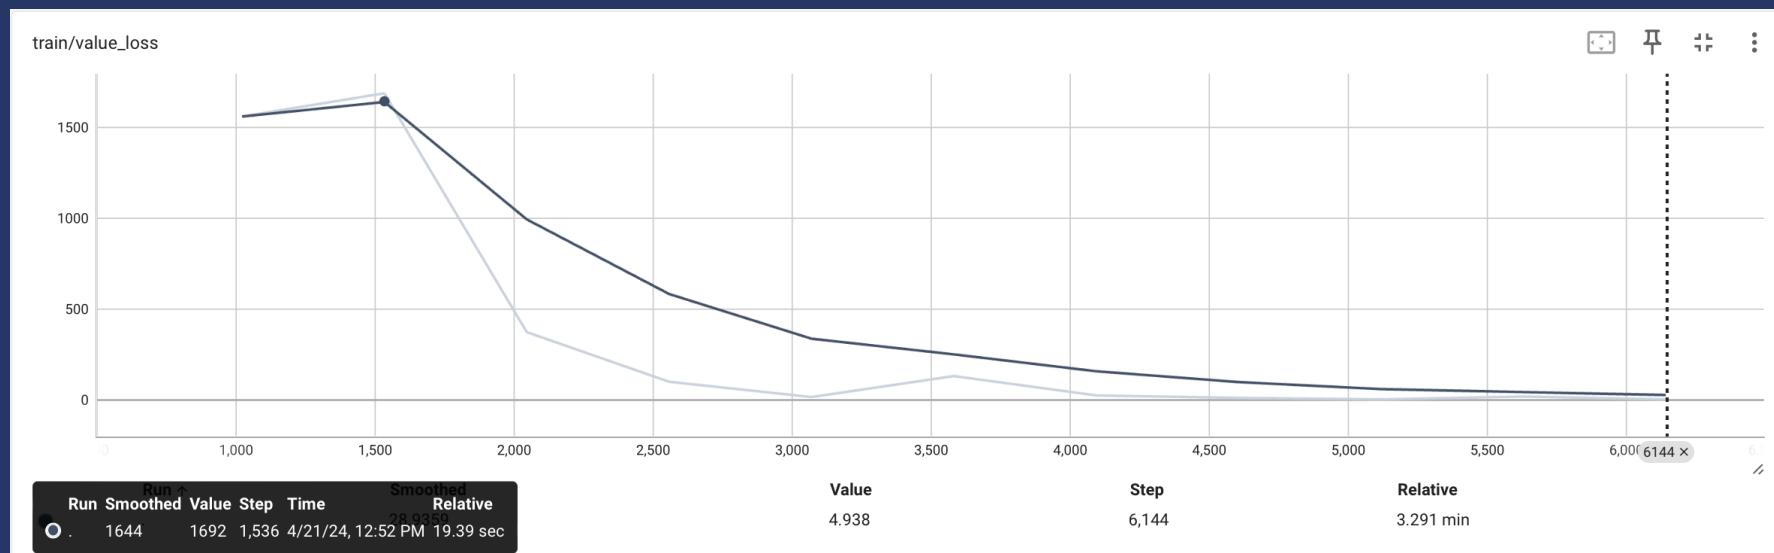
|                      |           |
|----------------------|-----------|
| rollout/             |           |
| ep_len_mean          | 27        |
| ep_rew_mean          | 252       |
| time/                |           |
| fps                  | 24        |
| iterations           | 554       |
| time_elapsed         | 11655     |
| total_timesteps      | 283648    |
| train/               |           |
| approx_kl            | 0.0       |
| clip_fraction        | 0         |
| clip_range           | 0.2       |
| entropy_loss         | -4.51e-06 |
| explained_variance   | 1         |
| learning_rate        | 0.001     |
| loss                 | 0.0231    |
| n_updates            | 5530      |
| policy_gradient_loss | 1.96e-09  |
| value_loss           | 0.0934    |



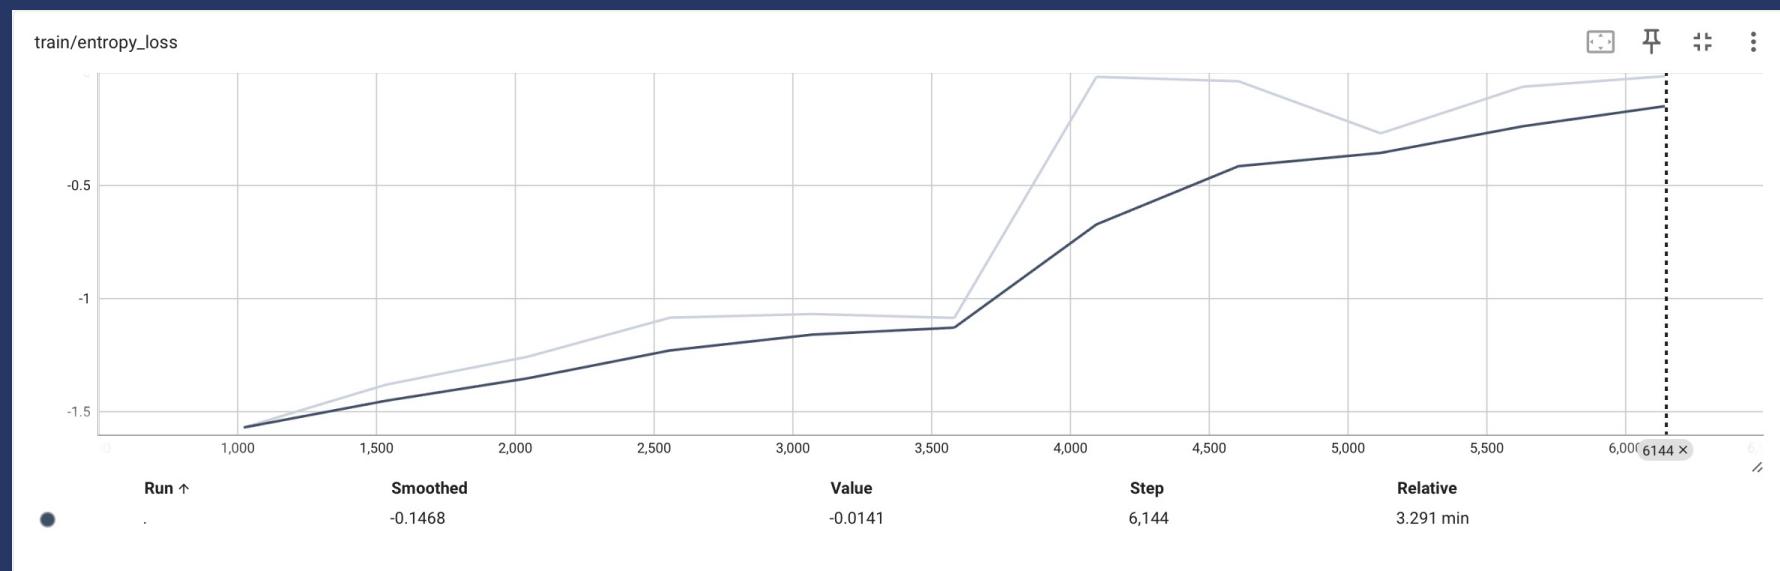
# Let's Talk Graphs – Episode Rewards



# Let's Talk Graphs – Value Loss



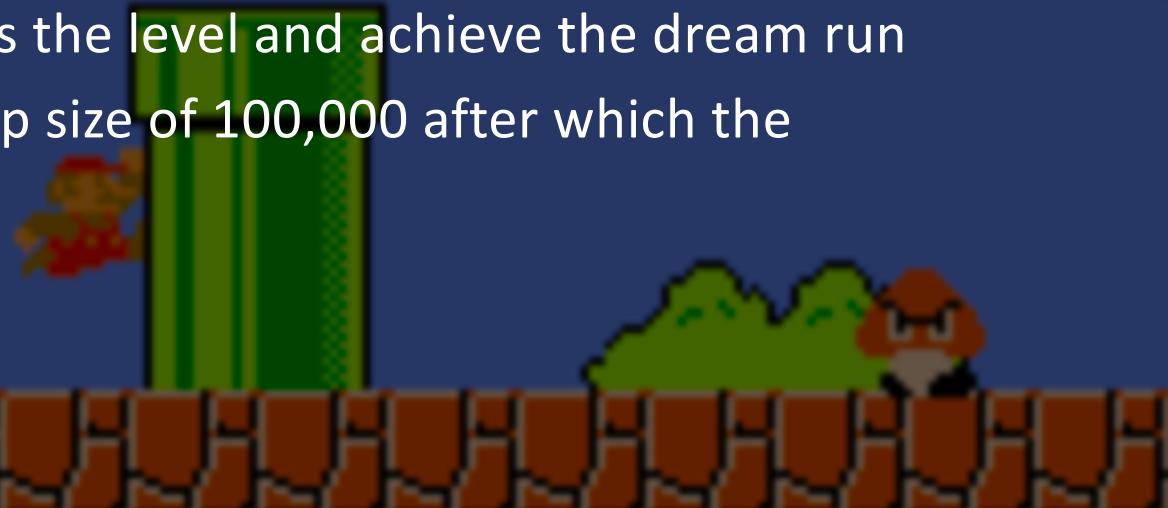
# Let's Talk Graphs – Entropy Loss



# Challenges

---

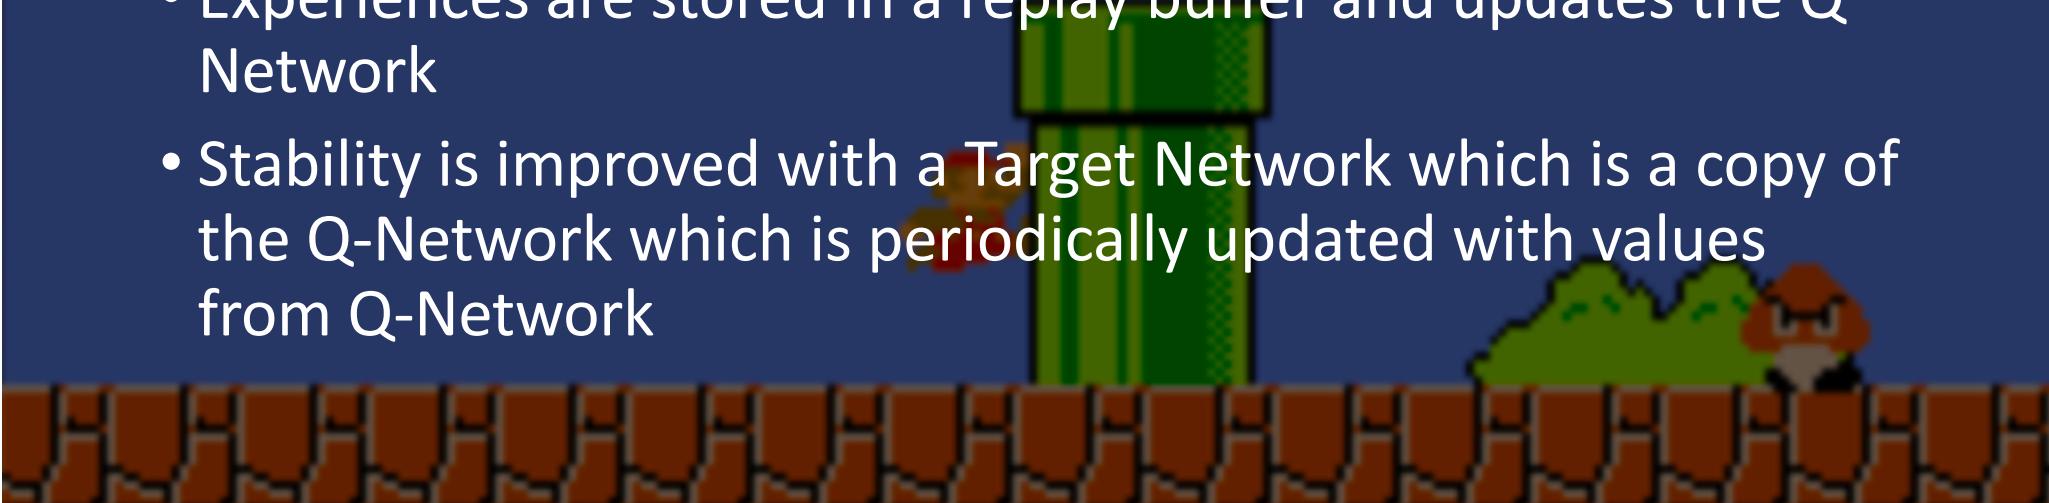
- Not having enough computation power and resources has been a big challenge for training this algorithm
- As per our research, a total step size of 4 million is required for the PPO algorithm to converge at an optimal stage
- That is where Mario can pass the level and achieve the dream run
- I was able to have a total step size of 100,000 after which the algorithm breaks



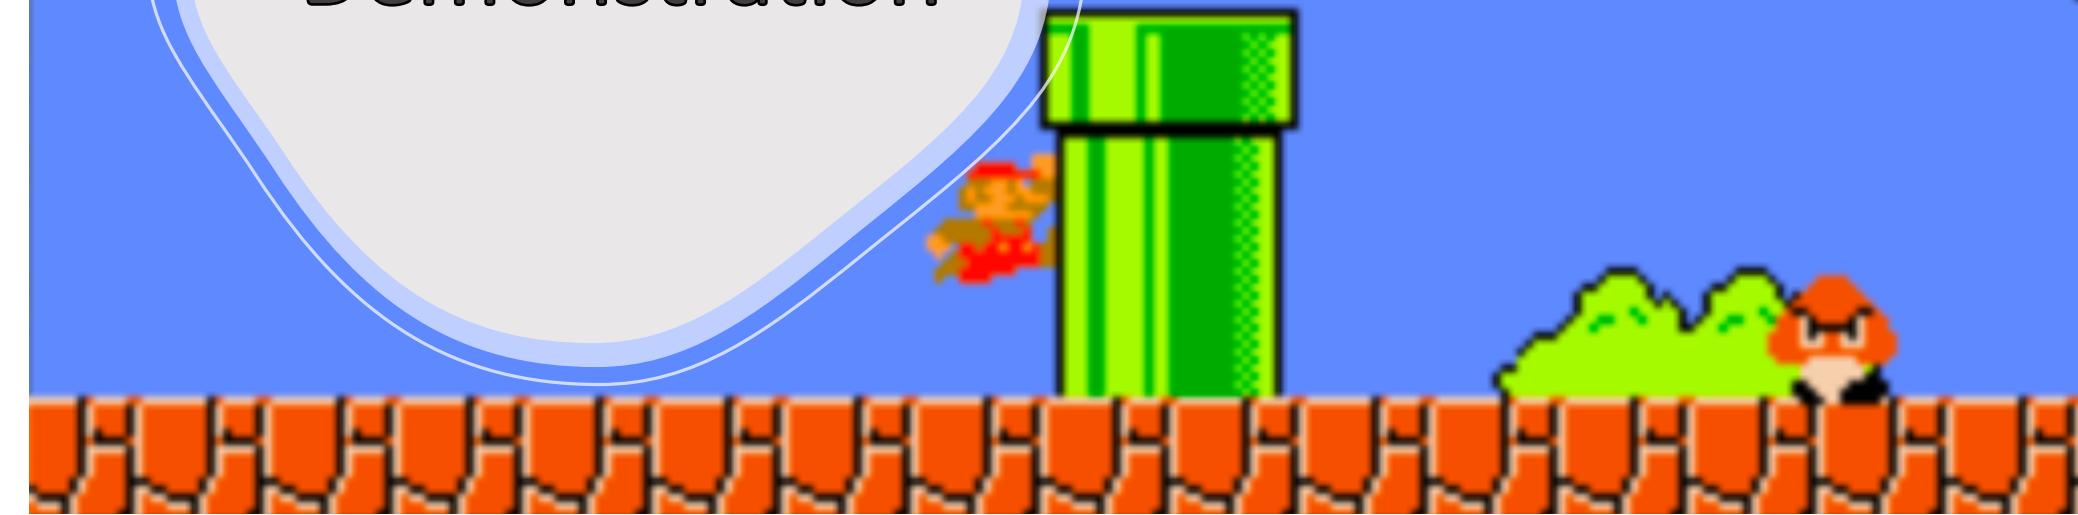
# Deep Q-Network (DQN)

---

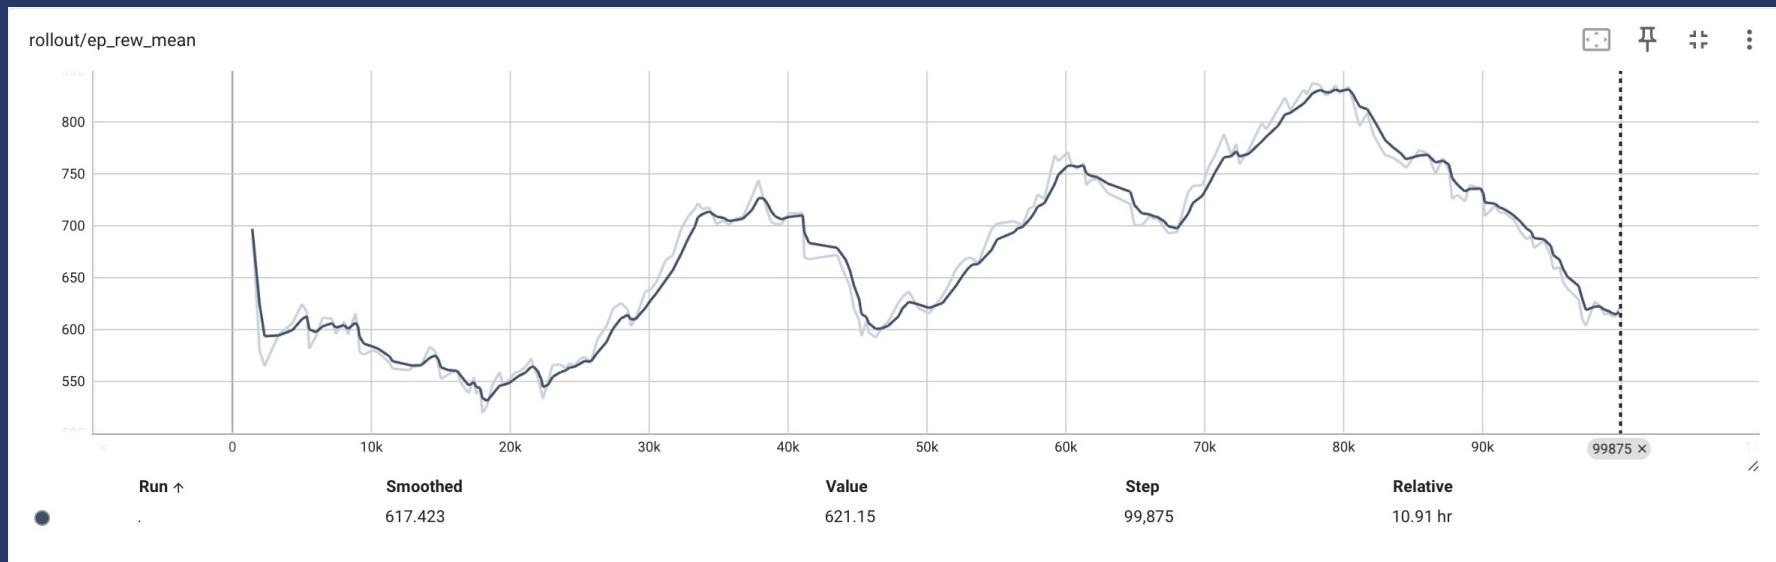
- A variant of Q Learning algorithm popular with Reinforcement Learning for game agents
- A CNN is used to approximate the Q function
- Experiences are stored in a replay buffer and updates the Q Network
- Stability is improved with a Target Network which is a copy of the Q-Network which is periodically updated with values from Q-Network



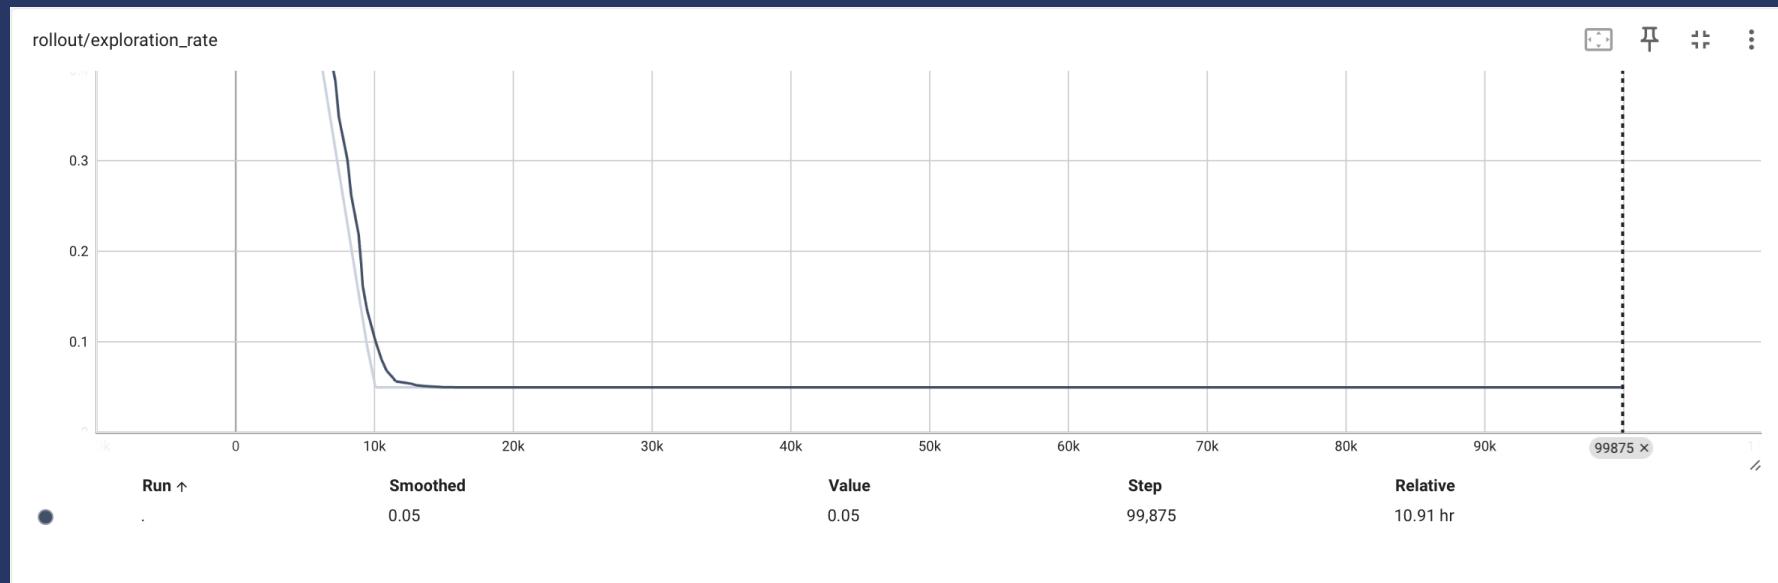
# Code Demonstration



# Let's Talk Graphs – Episode Rewards



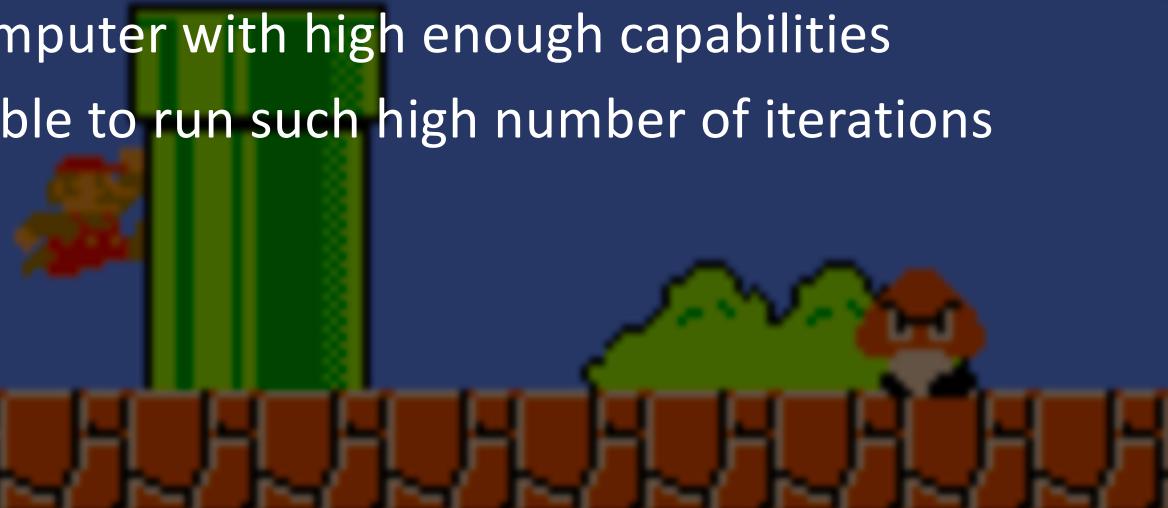
# Let's Talk Graphs – Exploration Rate



# Challenges

---

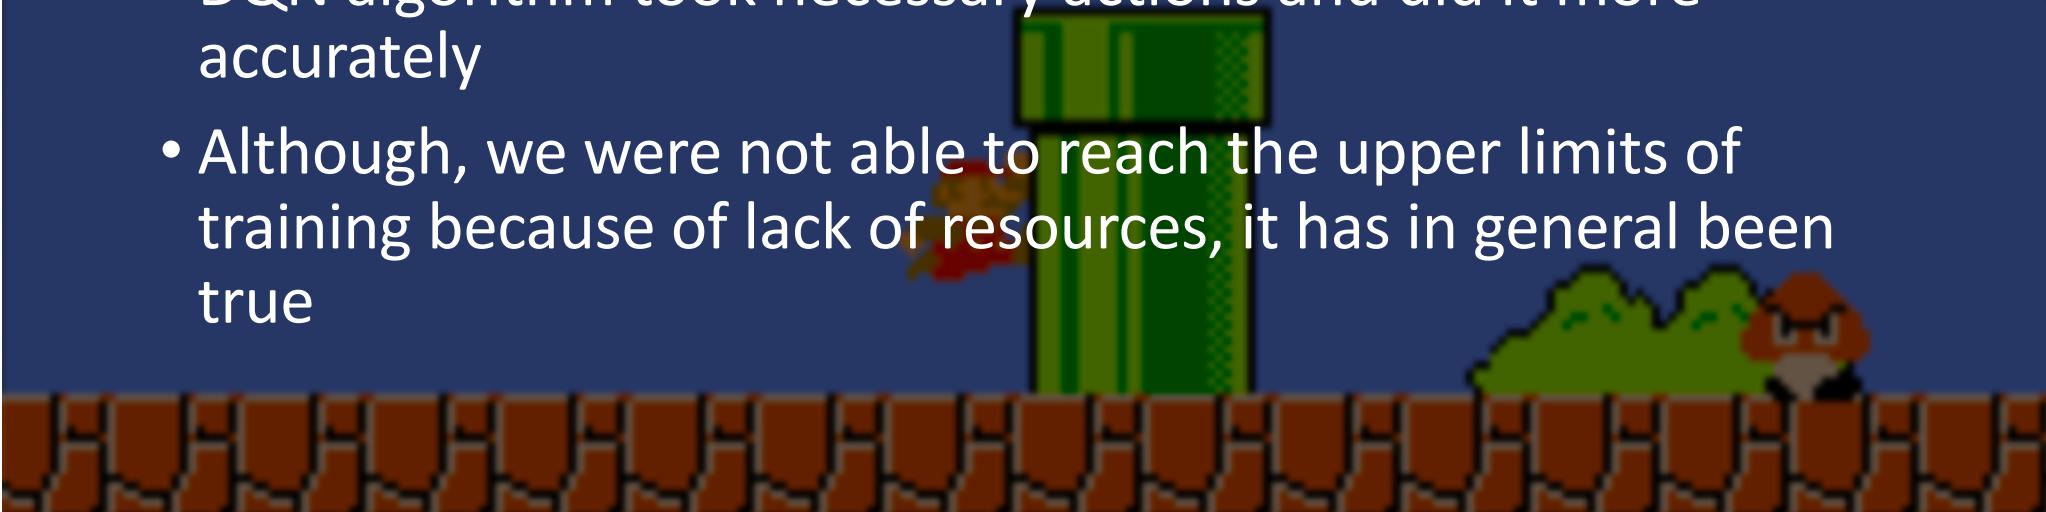
- Similar challenges compared to PPO
- Computation power and resources that we have are not enough to run the training model
- Model converges to the optimal value after 50,000 iterations which will take days to run on a computer with high enough capabilities
- Our school laptops are not able to run such high number of iterations
- RAM is not enough



# Conclusion - Lessons Learned – PPO vs DQN

---

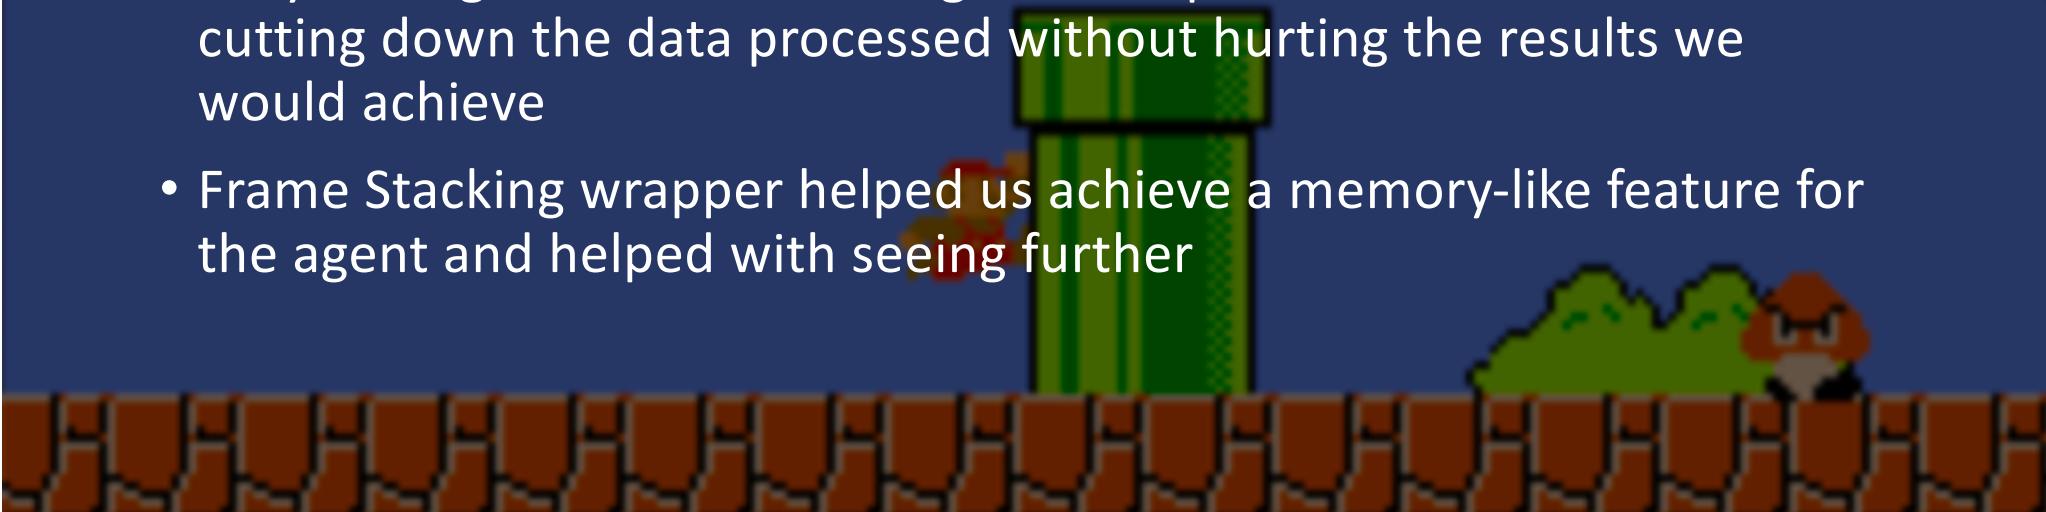
- Theoretically DQN is supposed to perform better than PPO, and it seems like through our exploration this has been the case
- DQN algorithm took necessary actions and did it more accurately
- Although, we were not able to reach the upper limits of training because of lack of resources, it has in general been true



# Conclusion – Wrapper Strategies

---

- Wrapping the environment and the agent within a few strategies has been very interesting
- We reduced the computational load for by about a half!
- Gray scaling and Frame Resizing have helped us achieve this idea of cutting down the data processed without hurting the results we would achieve
- Frame Stacking wrapper helped us achieve a memory-like feature for the agent and helped with seeing further



# Conclusion – Overall

---

- We learned about numerous concept of reinforcement learning and how to evaluate rewards
- The reward function for Mario is very complex, considering how many different thing we need to consider
- It was designed by the writers of PPO and DQN respectively and the software for the reward function is owned by them
- Both the algorithms do a very good job at trying achieve convergence, DQN being the overall better, but this decision is heavily based on the action space chosen and the description of the environment

